

What Is TL?

- Introduction & Motivation
- Adapting Neural Networks
- Process

Transfer Learning

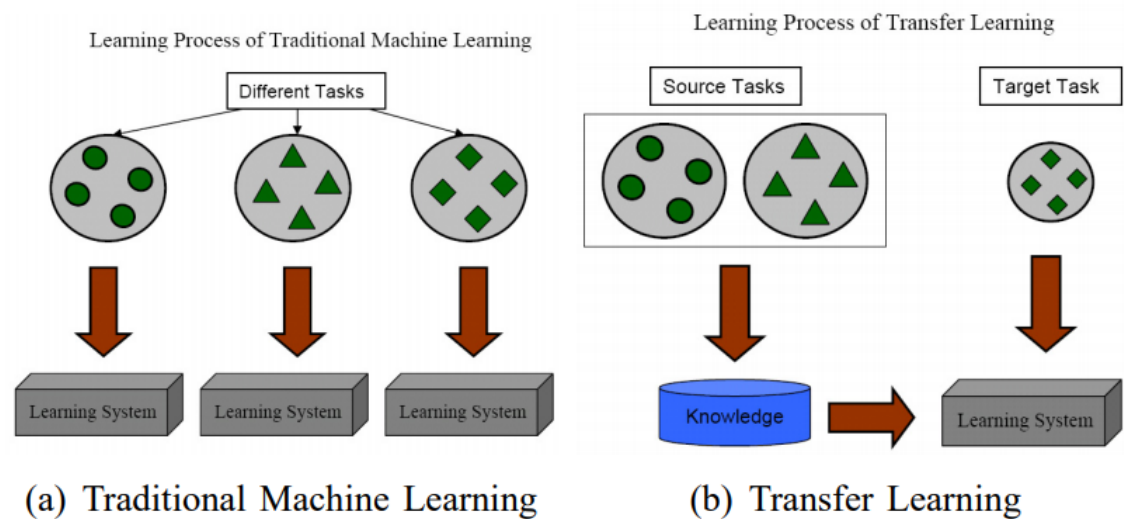
Transferring the knowledge of one model to perform a new task.

"Domain Adaptation"

Motivation

- Lots of data, time, resources needed to train and tune a neural network from scratch
 - An ImageNet deep neural net can take weeks to train and fine-tune from scratch.
 - Unless you have 256 GPUs, possible to achieve in 1 hour
(<https://research.fb.com/publications/accurate-large-minibatch-sgd-training-imagenet-in-1-hour/>).
- Cheaper, faster way of adapting a neural network by exploiting their generalization properties

Traditional vs. Transfer Learning



(image: [Survey on Transfer Learning](#)

(<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.147.9185&rep=rep1&type=pdf>))

Transfer Learning Types

Type	Description	Examples
Inductive	Adapt existing supervised training model on new labeled dataset	Classification, Regression
Transductive	Adapt existing supervised training model on new unlabeled dataset	Classification, Regression
Unsupervised	Adapt existing unsupervised training model on new unlabeled dataset	Clustering, Dimensionality Reduction

Transfer Learning Applications

- Image classification (most common): learn new image classes
- Text sentiment classification
- Text translation to new languages
- Speaker adaptation in speech recognition
- Question answering

Transfer Learning Services

Transfer learning is used in many "train your own AI model" services:

- just upload 5-10 images to train a new model! in minutes!

Easily customize your own state-of-the-art computer vision models for your unique use case. Just upload a few labeled images and let Custom Vision Service do the hard work. With just one click, you can export trained models to be run on device or as Docker containers.



(image: <https://azure.microsoft.com/en-us/services/cognitive-services/custom-vision-service/>
(<https://azure.microsoft.com/en-us/services/cognitive-services/custom-vision-service/>))


Transfer Learning in Neural Networks


Neural Network Layers: General to Specific

- Bottom/first/earlier layers: general learners
 - Low-level notions of edges, visual shapes
- Top/last/later layers: specific learners
 - High-level features such as eyes, feathers


Note: the top/bottom notation is confusing, I'd avoid it

Example: VGG 16 Filters







conv3_1: a few of the 256 filters



conv4_1: a few of the 512 filters




conv5_1: a few of the 512 filters



<https://blog.keras.io/how-convolutional-neural-networks-see-the-world.html> (<https://blog.keras.io/how-convolutional-neural-networks-see-the-world.html>)

Transfer Learning Overview

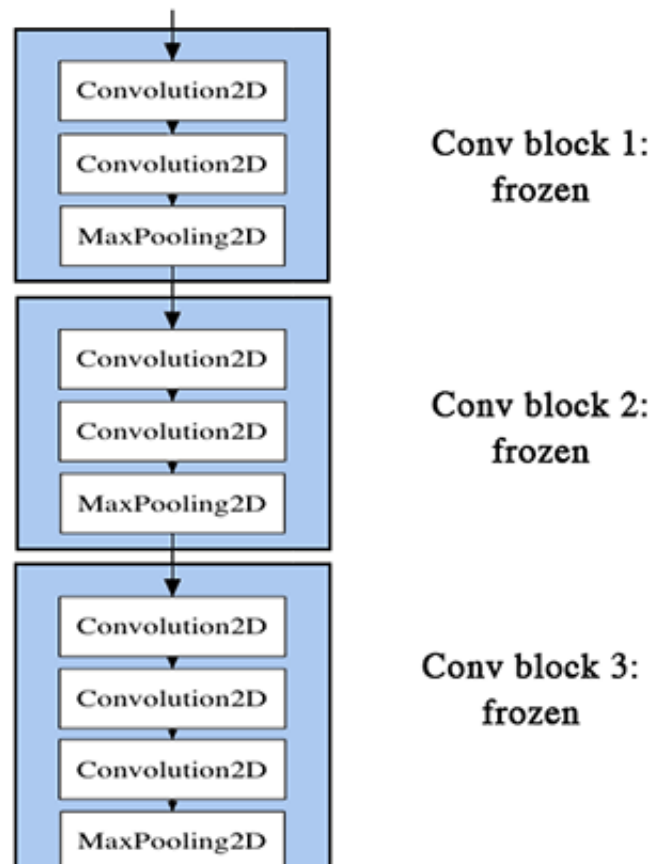


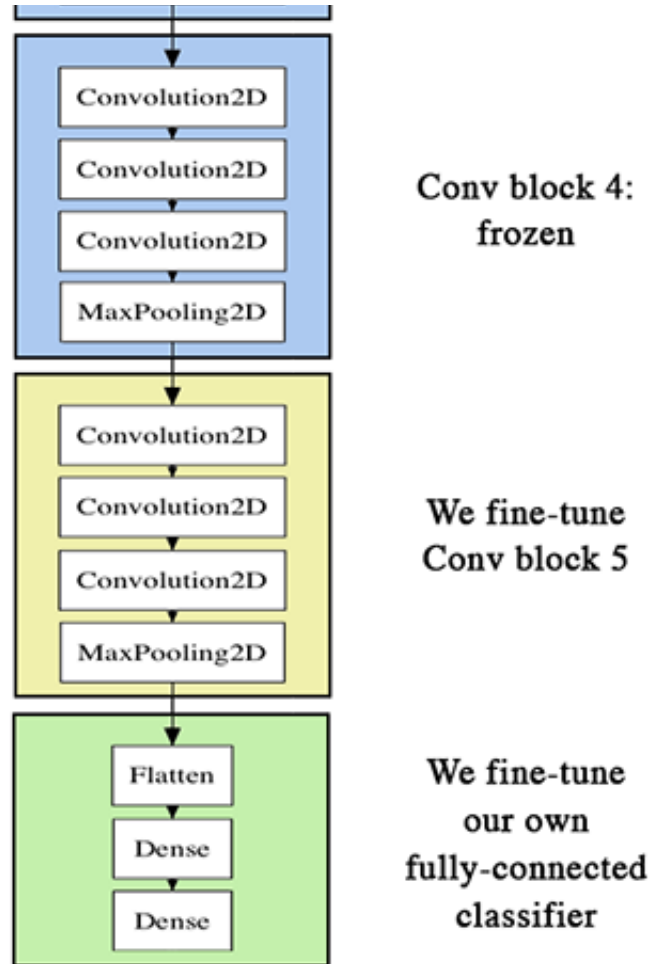
(image: [Aghamirzaie & Salomon](#)
(<http://slideplayer.com/slide/8370683/>))

Process

1. Start with pre-trained network
2. Partition network into:
 - Featurizers: identify which layers to keep
 - Classifiers: identify which layers to replace
3. Re-train classifier layers with new data
4. Unfreeze weights and fine-tune whole network with smaller learning rate

Freezing and Fine-tuning






(image: <http://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>

(<http://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>))

Which layers to re-train?

- Depends on the domain
- Start by re-training the last layers (last full-connected and last convolutional)
 - work backwards if performance is not satisfactory

Example



(image: <http://arxiv.org/abs/1411.1792>)

(<http://arxiv.org/abs/1411.1792>)

When and how to fine-tune?

Suppose we have model A, trained on dataset A Q: How do we apply transfer learning to dataset B to create model B?

Dataset size	Dataset similarity	Recommendation
Large	Very different	Train model B from scratch, initialize weights from model A
Large	Similar	OK to fine-tune (less likely to overfit)
Small	Very different	Train classifier using the earlier layers (later layers won't help much)
Small	Similar	Don't fine-tune (overfitting). Train a linear classifier

<https://cs231n.github.io/transfer-learning/>
(<https://cs231n.github.io/transfer-learning/>)

Learning Rates

- Training linear classifier: typical learning rate
- Fine-tuning: use smaller learning rate to avoid distorting the existing weights
 - Assumes weights are close to "good"